

Mapping and spatial analysis with ACS data in R

SSDAN Workshop Series

Kyle Walker

February 15, 2023

About me

- Associate Professor of Geography at TCU
- Spatial data science researcher and consultant
- Package developer: **tidycensus**, **tigris**, **mapboxapi**, **crsuggest**, **idbr** (R), **pygris** (Python)
- Book: [*Analyzing US Census Data: Methods, Maps and Models in R*](#)
 - Print release date February 16 (tomorrow!)
 - To support these workshops: [buy on Amazon](#) or [direct from CRC Press](#)



SSDAN workshop series

- Last week (February 8): [Working with the 2021 American Community Survey with R and tidycensus](#)
- Today: Mapping and spatial analysis with ACS data in R
- Next week (February 22): Working with geographic data and making maps in Python

Today's agenda

- Hour 1: Working with "spatial" American Community Survey data
- Hour 2: Making maps of ACS data in R
- Hour 3: Applications: segregation, diversity, and location intelligence

Part 1: Working with "spatial" American Community Survey data

The American Community Survey

- Annual survey of 3.5 million US households
- Covers topics not available in decennial US Census data (e.g. income, education, language, housing characteristics)
- Available as 1-year estimates (for geographies of population 65,000 and greater) and 5-year estimates (for geographies down to the block group)
 - 2020 1-year data only available as [experimental estimates](#)
- Data delivered as *estimates* characterized by *margins of error*

How to get ACS data

- data.census.gov is the main, revamped interactive data portal for browsing and downloading Census datasets, including the ACS
- [The US Census Application Programming Interface \(API\)](#) allows developers to access Census data resources programmatically

tidycensus

- R interface to the Decennial Census, American Community Survey, Population Estimates Program, and Public Use Microdata Series APIs
- Key features:
 - Wrangles Census data internally to return tidyverse-ready format (or traditional wide format if requested);
 - Automatically downloads and merges Census geometries to data for mapping (next week's workshop!);
 - Includes tools for handling margins of error in the ACS and working with survey weights in the ACS PUMS;
 - States and counties can be requested by name (no more looking up FIPS codes!)

R and RStudio

- R: programming language and software environment for data analysis (and wherever else your imagination can take you!)
- RStudio: integrated development environment (IDE) for R developed by [Posit](#)
- Posit Cloud: run RStudio with today's workshop pre-configured at <https://posit.cloud/content/5377428>

Getting started with tidycensus

- To get started, install the packages you'll need for today's workshop
- If you are using the Posit Cloud environment, these packages are already installed for you

```
install.packages(c("tidycensus", "tidyverse"))
```

- Optional, to run advanced examples:

```
install.packages(c("mapview", "mapedit", "mapboxapi",  
                  "leafsync", "spdep", "segregation",  
                  "ggiraph"))
```

Optional: your Census API key

- tidycensus (and the Census API) can be used without an API key, but you will be limited to 500 queries per day
- Power users: visit https://api.census.gov/data/key_signup.html to request a key, then activate the key from the link in your email.
- Once activated, use the `census_api_key()` function to set your key as an environment variable

```
library(tidycensus)

census_api_key("YOUR KEY GOES HERE", install = TRUE)
```

Spatial Census data in tidycensus

Spatial Census data: the old way

Traditionally, getting "spatial" Census data requires:

- Fetching shapefiles from the Census website;
- Downloading a CSV of data, cleaning/formatting it;
- Loading geometries and data into your GIS of choice;
- Aligning key fields in your GIS and joining your data

Spatial Census data with `get_acs()`

- The `get_acs()` function is your portal to access ACS data using tidycensus
- The two required arguments are `geography` and `variables`. The function defaults to the 2017-2021 5-year ACS
- The argument `geometry = TRUE` returns pre-joined geometry along with your ACS data!

```
library(tidycensus)

texas_income <- get_acs(
  geography = "county",
  variables = "B19013_001",
  state = "TX",
  year = 2021,
  geometry = TRUE
)

plot(texas_income[ "estimate" ])
```



Looking under the hood: *simple features* in R



- The sf package implements a *simple features data model* for vector spatial data in R
- Vector geometries: *points*, *lines*, and *polygons* stored in a list-column of a data frame

- Spatial data are returned with five data columns: GEOID, NAME, variable, estimate, and moe, along with a geometry column representing the shapes of locations

```
texas_income
```

```
## Simple feature collection with 254 features and 5 fields
## Geometry type: MULTIPOLYGON
## Dimension:      XY
## Bounding box:   xmin: -106.6456 ymin: 25.83738 xmax: -93.50829 ymax: 36.5007
## Geodetic CRS:   NAD83
## First 10 features:
##      GEOID      NAME      variable estimate moe
## 1  48053  Burnet County, Texas B19013_001   65363 4694
## 2  48057  Calhoun County, Texas B19013_001   61887 9517
## 3  48341    Moore County, Texas B19013_001   55543 3786
## 4  48185    Grimes County, Texas B19013_001   59086 4414
## 5  48035    Bosque County, Texas B19013_001   59328 3702
## 6  48471    Walker County, Texas B19013_001   44104 1820
## 7  48407 San Jacinto County, Texas B19013_001   46678 6190
## 8  48005   Angelina County, Texas B19013_001   52377 3075
## 9  48443   Terrell County, Texas B19013_001   47012 5245
## 10 48421   Sherman County, Texas B19013_001   55667 9007
##      geometry
## 1  MULTIPOLYGON (((-98.45924 3...
## 2  MULTIPOLYGON (((-96.77168 2...
## 3  MULTIPOLYGON (((-102.163 36...
## 4  MULTIPOLYGON (((-96.18831 3...
## 5  MULTIPOLYGON (((-98.00068 3...
```

```
## 6 MULTIPOLYGON (((-95.86227 3...
```

Interactive viewing with `mapview()`



- The **mapview** package allows for interactive viewing of spatial data in R

```
library(mapview)
```

```
mapview(texas_income, zcol = "estimate")
```



Understanding geography and variables in tidycensus

US Census Geography



Source: [US Census Bureau](#)

Geography in tidycensus

- Information on available geographies, and how to specify them, can be found [in the tidycensus documentation](#)



Searching for variables

- To search for variables, use the `load_variables()` function along with a year and dataset
- For the 2021 5-year ACS, use "acs5" for the Detailed Tables; "acs5/profile" for the Data Profile; "acs5/subject" for the Subject Tables; and "acs5/cprofile" for the Comparison Profile
- The `View()` function in RStudio allows for interactive browsing and filtering

```
vars <- load_variables(2021, "acs5")  
View(vars)
```

Small-area spatial demographic data

- Smaller areas like Census tracts or block groups are available with `geography = "tract"` or `geography = "block group"`; a county can optionally be specified to focus your query

```
king_income <- get_acs(  
  geography = "tract",  
  variables = "B19013_001",  
  state = "WA",  
  county = "King",  
  geometry = TRUE  
)  
  
mapview(king_income, zcol = "estimate")
```




Spatial data structure in tidycensus

"Tidy" or long-form data

- The default data structure returned by **tidycensus** is "tidy" or long-form data, with variables by geography stacked by row
- For spatial data, this means that geometries will also be stacked, which is helpful for group-wise analysis and visualization

```
orange_race <- get_acs(  
  geography = "tract",  
  variables = c(  
    Hispanic = "DP05_0071P",  
    White = "DP05_0077P",  
    Black = "DP05_0078P",  
    Asian = "DP05_0080P"  
  ),  
  state = "CA",  
  county = "Orange",  
  geometry = TRUE  
)
```

```
orange_race
```

```
## Simple feature collection with 2456 features and 5 fields (with 4 geometries empty)
## Geometry type: MULTIPOLYGON
## Dimension:      XY
## Bounding box:   xmin: -118.1154 ymin: 33.38779 xmax: -117.4133 ymax: 33.94764
## Geodetic CRS:  NAD83
## First 10 features:
##      GEOID      NAME variable estimate
## 1  06059011402 Census Tract 114.02, Orange County, California Hispanic      21.4
## 2  06059011402 Census Tract 114.02, Orange County, California   White      51.7
## 3  06059011402 Census Tract 114.02, Orange County, California   Black       1.3
## 4  06059011402 Census Tract 114.02, Orange County, California   Asian      19.6
## 5  06059087403 Census Tract 874.03, Orange County, California Hispanic      87.4
## 6  06059087403 Census Tract 874.03, Orange County, California   White       8.1
## 7  06059087403 Census Tract 874.03, Orange County, California   Black       0.3
## 8  06059087403 Census Tract 874.03, Orange County, California   Asian       3.8
## 9  06059011716 Census Tract 117.16, Orange County, California Hispanic      27.7
## 10 06059011716 Census Tract 117.16, Orange County, California   White      38.6
##      moe      geometry
## 1  5.3 MULTIPOLYGON (((-117.9137 3...
## 2  8.1 MULTIPOLYGON (((-117.9137 3...
## 3  1.9 MULTIPOLYGON (((-117.9137 3...
## 4  9.2 MULTIPOLYGON (((-117.9137 3...
## 5  5.4 MULTIPOLYGON (((-117.9154 3...
## 6  3.6 MULTIPOLYGON (((-117.9154 3...
## 7  0.4 MULTIPOLYGON (((-117.9154 3...
## 8  3.3 MULTIPOLYGON (((-117.9154 3...
```

"Wide" data

- The argument `output = "wide"` spreads Census variables across the columns, returning one row per geographic unit and one column per variable
- This will be a more familiar data structure for traditional desktop GIS users

```
orange_race_wide <- get_acs(  
  geography = "tract",  
  variables = c(  
    Hispanic = "DP05_0071P",  
    White = "DP05_0077P",  
    Black = "DP05_0078P",  
    Asian = "DP05_0080P"  
  ),  
  state = "CA",  
  county = "Orange",  
  geometry = TRUE,  
  output = "wide"  
)
```

```
orange_race_wide
```

```
## Simple feature collection with 614 features and 10 fields (with 1 geometry empty)
## Geometry type: MULTIPOLYGON
## Dimension:      XY
## Bounding box:   xmin: -118.1154 ymin: 33.38779 xmax: -117.4133 ymax: 33.94764
## Geodetic CRS:   NAD83
## First 10 features:
```

##	GEOID	NAME	HispanicE
## 1	06059011402	Census Tract 114.02, Orange County, California	21.4
## 2	06059087403	Census Tract 874.03, Orange County, California	87.4
## 3	06059011716	Census Tract 117.16, Orange County, California	27.7
## 4	06059075504	Census Tract 755.04, Orange County, California	37.6
## 5	06059063902	Census Tract 639.02, Orange County, California	25.7
## 6	06059099222	Census Tract 992.22, Orange County, California	15.8
## 7	06059088403	Census Tract 884.03, Orange County, California	57.8
## 8	06059110108	Census Tract 1101.08, Orange County, California	39.2
## 9	06059011403	Census Tract 114.03, Orange County, California	56.6
## 10	06059087505	Census Tract 875.05, Orange County, California	68.9

```
## HispanicM WhiteE WhiteM BlackE BlackM AsianE AsianM
## 1      5.3    51.7    8.1    1.3    1.9    19.6    9.2
## 2      5.4     8.1     3.6     0.3     0.4     3.8     3.3
## 3      7.2    38.6     9.5     0.5     0.9    30.4     5.2
## 4      9.4    46.1     7.1     1.9     2.2     9.5     3.5
## 5      7.1    56.6     8.7     1.7     1.3    10.4     4.9
## 6      5.7    18.6     8.9     1.4     1.6    58.5     9.6
## 7      9.5    16.3     5.9     1.1     0.8    24.2     7.9
## 8     15.7    40.4    11.4     3.4     3.5    12.4     6.2
```

Working with Census geometry

Census geometry and the **tigris** R package



- tidycensus uses the **tigris** R package internally to acquire Census shapefiles

Problem: interior water areas

- Let's re-visit the King County income-by-Census tract map
- Areas like Mercer Island are obscured as Census tracts cover the entirety of Lake Washington
- `erase_water()` in the **tigris** package offers a solution; it automates the removal of water areas from your shapes

```
library(tigris)
library(sf)
sf_use_s2(FALSE)

king_erase <- erase_water(king_income,
                          area_threshold = 0.9,
                          year = 2021)

mapview(king_erase, zcol = "estimate")
```



Part 1 exercises

1. Use the `load_variables()` function to find a variable that interests you that we haven't used yet.
2. Use `get_acs()` to fetch spatial ACS data on that variable for a geography and location of your choice, then use `mapview()` to display your data interactively.

Part 2: Mapping ACS data

Mapping in R

- R has a robust set of tools for cartographic visualization that make it a suitable alternative to desktop GIS software in many instances
- Popular packages for cartography include **ggplot2**, **tmap**, and **maps**
- Today, we'll be focusing on ggplot2; [see Chapter 6 of my book for similar examples using tmap](#)

ggplot2 and `geom_sf()`

- ggplot2: R's most popular visualization package (over 105 million downloads!)
- ggplot2 graphics are defined by an *aesthetic mapping* and one or more *geoms*
- `geom_sf()` is a special geom that interprets the geometry type of your spatial data and visualizes it accordingly
- As a result, we can make attractive maps using familiar ggplot2 syntax

Mapping ACS data with ggplot2

Continuous choropleth

- By default, ggplot2 will apply a continuous color palette to create **choropleth** maps
- Choropleth maps: the shading of a polygon / shape is mapped to a data attribute

```
library(tidyverse)

orange_hispanic <- filter(orange_race, variable == "Hispanic")

ggplot(orange_hispanic, aes(fill = estimate)) +
  geom_sf()
```




Continuous choropleth with styling

- We can apply some styling to customize our choropleth maps
- Used here: a [viridis](#) color palette, which is built-in to ggplot2

```
ggplot(orange_hispanic, aes(fill = estimate)) +  
  geom_sf() +  
  theme_void() +  
  scale_fill_viridis_c(option = "rocket") +  
  labs(title = "Percent Hispanic by Census tract",  
        subtitle = "Orange County, California",  
        fill = "ACS estimate",  
        caption = "2017-2021 ACS | tidycensus R package")
```



Classed choropleth

- We can also create a binned choropleth; ggplot2 will identify "pretty" breaks, or custom breaks can be supplied

```
ggplot(orange_hispanic, aes(fill = estimate)) +  
  geom_sf() +  
  theme_void() +  
  scale_fill_viridis_b(option = "rocket", n.breaks = 6) +  
  labs(title = "Percent Hispanic by Census tract",  
        subtitle = "Orange County, California",  
        fill = "ACS estimate",  
        caption = "2017-2021 ACS | tidycensus R package")
```



Faceted choropleth maps

- Spatial ACS data in tidy (long) format can be *faceted* by a grouping variable, allowing for comparative mapping

```
ggplot(orange_race, aes(fill = estimate)) +  
  geom_sf(color = NA) +  
  theme_void() +  
  scale_fill_viridis_c(option = "rocket") +  
  facet_wrap(~variable) +  
  labs(title = "Race / ethnicity by Census tract",  
        subtitle = "Orange County, California",  
        fill = "ACS estimate (%)",  
        caption = "2017-2021 ACS | tidycensus R package")
```



Mapping count data

Mapping count data

- At times, you'll want to show variations in *counts* rather than rates on your maps of ACS data
- Choropleth maps are poorly suited for count data
- Let's grab some count data for race / ethnicity and consider some alternatives

```
orange_race_counts <- get_acs(  
  geography = "tract",  
  variables = c(  
    Hispanic = "DP05_0071",  
    White = "DP05_0077",  
    Black = "DP05_0078",  
    Asian = "DP05_0080"  
  ),  
  state = "CA",  
  county = "Orange",  
  geometry = TRUE  
)
```

Graduated symbol maps

- Graduated symbol maps show difference on a map with the size of symbols (often circles)
- They are better for count data than choropleth maps as the shapes are directly comparable (unlike differentially-sized polygons)
- We'll need to convert our data to *centroids* to plot graduated symbols in ggplot2

```
library(sf)

orange_black <- filter(
  orange_race_counts,
  variable == "Black"
)

centroids <- st_centroid(orange_black)
```

Graduated symbol maps

- We'll first plot a base layer of Census tracts, then a layer of graduated symbols on top
- Use `scale_size_area()` to plot *proportional symbols*

```
ggplot() +  
  geom_sf(data = orange_black, color = "black", fill = "lightgrey") +  
  geom_sf(data = centroids, aes(size = estimate),  
          alpha = 0.7, color = "navy") +  
  theme_void() +  
  labs(title = "Black population by Census tract",  
        subtitle = "2017-2021 ACS, Orange County, California",  
        size = "ACS estimate") +  
  scale_size_area(max_size = 6)
```



Dot-density mapping

- It can be difficult to show *heterogeneity* or *mixing* of different categories on maps
- Dot-density maps scatter dots proportionally to data size; dots can be colored to show mixing of categories
- Traditionally, dot-density maps are slow to make in R; tidycensus's `as_dot_density()` function addresses this

```
orange_race_dots <- as_dot_density(  
  orange_race_counts,  
  value = "estimate",  
  values_per_dot = 200,  
  group = "variable"  
)
```

```
orange_race_dots
```

```
## Simple feature collection with 15199 features and 5 fields
## Geometry type: POINT
## Dimension:      XY
## Bounding box:   xmin: -118.1149 ymin: 33.39638 xmax: -117.4263 ymax: 33.94577
## Geodetic CRS:  NAD83
## First 10 features:
```

##	GEOID	NAME	variable	estimate
## 1	06059086303	Census Tract 863.03, Orange County, California	White	2329
## 2	06059088301	Census Tract 883.01, Orange County, California	Asian	2519
## 3	06059099601	Census Tract 996.01, Orange County, California	Hispanic	2188
## 4	06059063906	Census Tract 639.06, Orange County, California	White	2306
## 5	06059062657	Census Tract 626.57, Orange County, California	White	2802
## 6	06059001202	Census Tract 12.02, Orange County, California	Hispanic	3059
## 7	06059075404	Census Tract 754.04, Orange County, California	White	1793
## 8	06059074103	Census Tract 741.03, Orange County, California	Hispanic	5255
## 9	06059088302	Census Tract 883.02, Orange County, California	Hispanic	1648
## 10	06059099239	Census Tract 992.39, Orange County, California	White	2674

```
## moe geometry
## 1 580 POINT (-117.8948 33.81615)
## 2 542 POINT (-117.9413 33.78907)
## 3 505 POINT (-118.0064 33.74013)
## 4 402 POINT (-117.9109 33.65461)
## 5 695 POINT (-117.8486 33.63091)
## 6 460 POINT (-117.9456 33.92972)
## 7 468 POINT (-117.8548 33.76288)
## 8 763 POINT (-117.8821 33.72321)
```

Dot-density mapping

- Like the graduated symbol map, we plot points over a base layer, but in this case with a much smaller size
- Use `override.aes` in `guide_legend()` to plot visible colors in the legend

```
ggplot() +  
  geom_sf(data = orange_black, color = "lightgrey", fill = "white") +  
  geom_sf(data = orange_race_dots, aes(color = variable), size = 0.01) +  
  scale_color_brewer(palette = "Set1") +  
  guides(color = guide_legend(override.aes = list(size = 3))) +  
  theme_void() +  
  labs(color = "Race / ethnicity",  
       caption = "2017-2021 ACS | 1 dot = approximately 200 people")
```



Interactive mapping and national US maps

Customizing interactive maps with `mapview()`

- `mapview()` accepts custom color palettes and labels, making it a suitable engine for interactive maps for presentations!

```
library(viridisLite)

colors <- rocket(n = 100)

mapview(orange_hispanic, zcol = "estimate",
        layer.name = "Percent Hispanic<br/>2017-2021 ACS",
        col.regions = colors)
```



Linked interactive maps with `mapview()`

- In `mapview`, layers can be stacked with the `+` operator or swiped between with the `|` operator
- **leafsync** takes this one step further by creating side-by-side synced maps

```
library(leafsync)

orange_white <- filter(orange_race, variable == "White")

m1 <- mapview(orange_hispanic, zcol = "estimate",
              layer.name = "Percent Hispanic<br/>2017-2021 ACS",
              col.regions = colors)

m2 <- mapview(orange_white, zcol = "estimate",
              layer.name = "Percent White<br/>2017-2021 ACS",
              col.regions = colors)

sync(m1, m2)
```

Problem: national US maps of the US

- These approaches to a common use-case -- mapping the entirety of the US -- doesn't work well. Let's take a look:

```
state_age <- get_acs(  
  geography = "state",  
  variables = "B01002_001",  
  year = 2021,  
  survey = "acs1",  
  geometry = TRUE  
)  
  
mapview(state_age, zcol = "estimate",  
  col.regions = plasma(7),  
  layer.name = "Median age<br/>2021 ACS")
```



Solution: shifting and rescaling US geometry

- The `shift_geometry()` function in the **tigris** package moves and optionally rescales Alaska, Hawaii, and Puerto Rico to help with national mapping
- `shift_geometry()` includes options to preserve relative position and/or area if desired

```
library(tigris)

age_shifted <- shift_geometry(state_age)

ggplot(age_shifted, aes(fill = estimate)) +
  geom_sf() +
  scale_fill_viridis_c(option = "plasma") +
  theme_void() +
  labs(fill = "Median age \n2021 ACS")
```



Interactivity with ggiraph

- **ggiraph**: Alternative approach for making **ggplot2** graphics interactive
- Includes *_interactive() versions of **ggplot2** geoms that can bring chart elements to life
- This includes geom_sf_interactive()

ggiraph example

```
library(ggiraph)

gg <- ggplot(age_shifted, aes(fill = estimate, data_id = GEOID,
                             tooltip = estimate)) +
  geom_sf_interactive() +
  scale_fill_viridis_c(option = "plasma") +
  theme_void() +
  labs(fill = "Median age\n2021 ACS")

girafe(ggobj = gg) %>%
  girafe_options(opts_hover(css = "fill:cyan;"))
```

Part 2 exercises

Use `load_variables(2021, "acs5/profile")` to find another percentage variable (ends in P) from the ACS Data Profile. Use that variable to try the following:

- Acquire spatial ACS data for your variable and make a customized choropleth map with **ggplot2**
- If time: try converting this map to an interactive map with **ggiraph** or **mapview**

Part 3: Applications: segregation, diversity, and spatial analysis

Analyzing segregation and diversity

The dissimilarity index

- The *dissimilarity index* is one of the most common measures used in the social sciences to calculate segregation
- The statistic reflects the percentage of a group's population that would have to move for *local percentages* to match the overall percentages of the area
- The index D is calculated as:

$$D = \frac{1}{2} \sum_{i=1}^N \left| \frac{a_i}{A} - \frac{b_i}{B} \right|$$

Calculating a dissimilarity index

- The **segregation** R package implements the dissimilarity index along with several other commonly-used indices for segregation and diversity
- It works very well with tidycensus and other tidy (long-form) demographic datasets as inputs

```
library(segregation)

orange_race_counts %>%
  filter(variable %in% c("White", "Hispanic")) %>%
  dissimilarity(
    group = "variable",
    unit = "GEOID",
    weight = "estimate"
  )
```

```
##      stat      est
## 1:      D 0.5192184
```

Group-wise segregation analysis

- Segregation analysis can also be used within tidyverse-style groupwise data analytic workflows
- Here, we'll grab some data for four counties; generate a `county` column for grouping; then use `group_modify()` to calculate groupwise dissimilarity indices (on the next slide)

```
la_race_counts <- get_acs(  
  geography = "tract",  
  variables = c(  
    Hispanic = "DP05_0071",  
    White = "DP05_0077",  
    Black = "DP05_0078",  
    Asian = "DP05_0080"  
  ),  
  state = "CA",  
  county = c("Orange", "Los Angeles",  
             "San Bernardino", "Riverside")  
) %>%  
  separate(NAME,  
    into = c("state", "county", "tract"),  
    sep = ", ")
```



```

la_race_counts %>%
  filter(variable %in% c("White", "Hispanic")) %>%
  group_by(county) %>%
  group_modify(
    ~dissimilarity(
      data = .x,
      group = "variable",
      unit = "GEOID",
      weight = "estimate"
    )
  )

```

```

## # A tibble: 4 × 3
## # Groups:   county [4]
##   county          stat    est
##   <chr>          <chr> <dbl>
## 1 Los Angeles County D      0.613
## 2 Orange County    D      0.519
## 3 Riverside County D      0.404
## 4 San Bernardino County D      0.426

```

Analyzing diversity: the entropy index

- As opposed to the dissimilarity index, the *entropy index* is a measure of evenness between multiple groups. It is computed as:

$$E = \sum_{r=1}^n Q_r \ln \frac{1}{Q_r}$$

- If the natural log is replaced with \log_k where k is the number of groups, the index will be maximized at 1

Calculating the entropy index

- The entropy index is implemented in the segregation package with **entropy**
- We'll need to use `group_modify()` to group by Census tract, and combine the results into an output dataset

```
orange_entropy <- orange_race_counts %>%  
  group_by(GEOID) %>%  
  group_modify(~tibble(  
    entropy = entropy(  
      data = .x,  
      group = "variable",  
      weight = "estimate",  
      base = 4  
    )  
  ))
```

Mapping the result

- Let's grab Census tracts directly using the **tigris** package to make a map

```
orange_tracts <- tracts("CA", "Orange", year = 2021, cb = TRUE)
orange_diversity_geo <- left_join(orange_tracts, orange_entropy, by = "GEOID")
```

```
ggplot(orange_diversity_geo, aes(fill = entropy)) +
  geom_sf() +
  scale_fill_viridis_c(option = "mako") +
  theme_void() +
  labs(fill = "Entropy index")
```



Spatial analysis with ACS data

Spatial analysis with ACS data

- **Spatial analysis:** the analysis of data in a way that takes into account, or directly focuses on, the data's spatial properties
- Spatial analysis is a huge topic; read [Chapter 7](#) and [Chapter 8](#) of my book for a wide variety of examples
- Today's examples: finding neighborhood hot-spots, analyzing demographics within proximity of a site

Spatial analysis: identifying hot-spots

- The main package for analyzing spatial relationships in R is **spdep**
- We will identify *spatial neighbors*; build a *spatial weights matrix* from those neighbors; then calculate a hot-spot statistic, the Getis-Ord Local G

```
library(spdep)

neighbors <- poly2nb(
  orange_diversity_geo,
  queen = TRUE
)

weights <- nb2listw(neighbors)

G <- localG(
  orange_diversity_geo$entropy,
  listw = weights
)
```


Hot-spot analysis

- As the local G values are z-scores, we can identify critical thresholds for "hot" or "cold" spots and map our results accordingly

```
orange_localG <- orange_diversity_geo %>%
  mutate(localG = G,
         Hotspot = case_when(
           localG >= 2.576 ~ "High cluster",
           localG <= -2.576 ~ "Low cluster",
           TRUE ~ "Not significant"
         ))

ggplot(orange_localG, aes(fill = Hotspot)) +
  geom_sf(color = "grey90") +
  scale_fill_manual(values = c("red", "blue", "grey")) +
  theme_void()
```



Problem: how to get data for a custom area?

- Census data are organized based on Census hierarchies: state, county, etc.
- What if you want data for a custom area that doesn't align with these hierarchies?
- Possible solution: `filter_by`

Demo: getting data for a custom shape

Demo: getting data around an address

Thank you, hope to see you next week!